

The background of the image is a vibrant, multi-colored gradient. It features broad diagonal bands of color, including shades of blue, purple, magenta, orange, and yellow, creating a dynamic and modern aesthetic. The AWS re:Invent logo is positioned on the left side of the image, rendered in white text.

AWS  
re:Invent

**SVS401 - R**

# Optimizing Your Serverless Applications

**Chris Munns**

Senior Manager/Principal Developer Advocate - Serverless  
Amazon Web Services

# About me

Chris Munns - [munns@amazon.com](mailto:munns@amazon.com), [@chrismunns](https://twitter.com/chrismunns)

- **Sr Manager/Principal Developer Advocate – Serverless**
- **New Yorker** (ehhh...ish.. kids/burbs/ya know?)

## Previously:

- AWS Business Development Manager – DevOps, July '15 - Feb '17
- AWS Solutions Architect Nov '11- Dec '14
- Formerly on operations teams @Etsy and @Meetup
- Little time at a hedge fund, Xerox, and a few other startups
- Rochester Institute of Technology: Applied Networking and Systems Administration '05
- Internet infrastructure geek



# Why are we here today?

Today's focus:




# Serverless applications

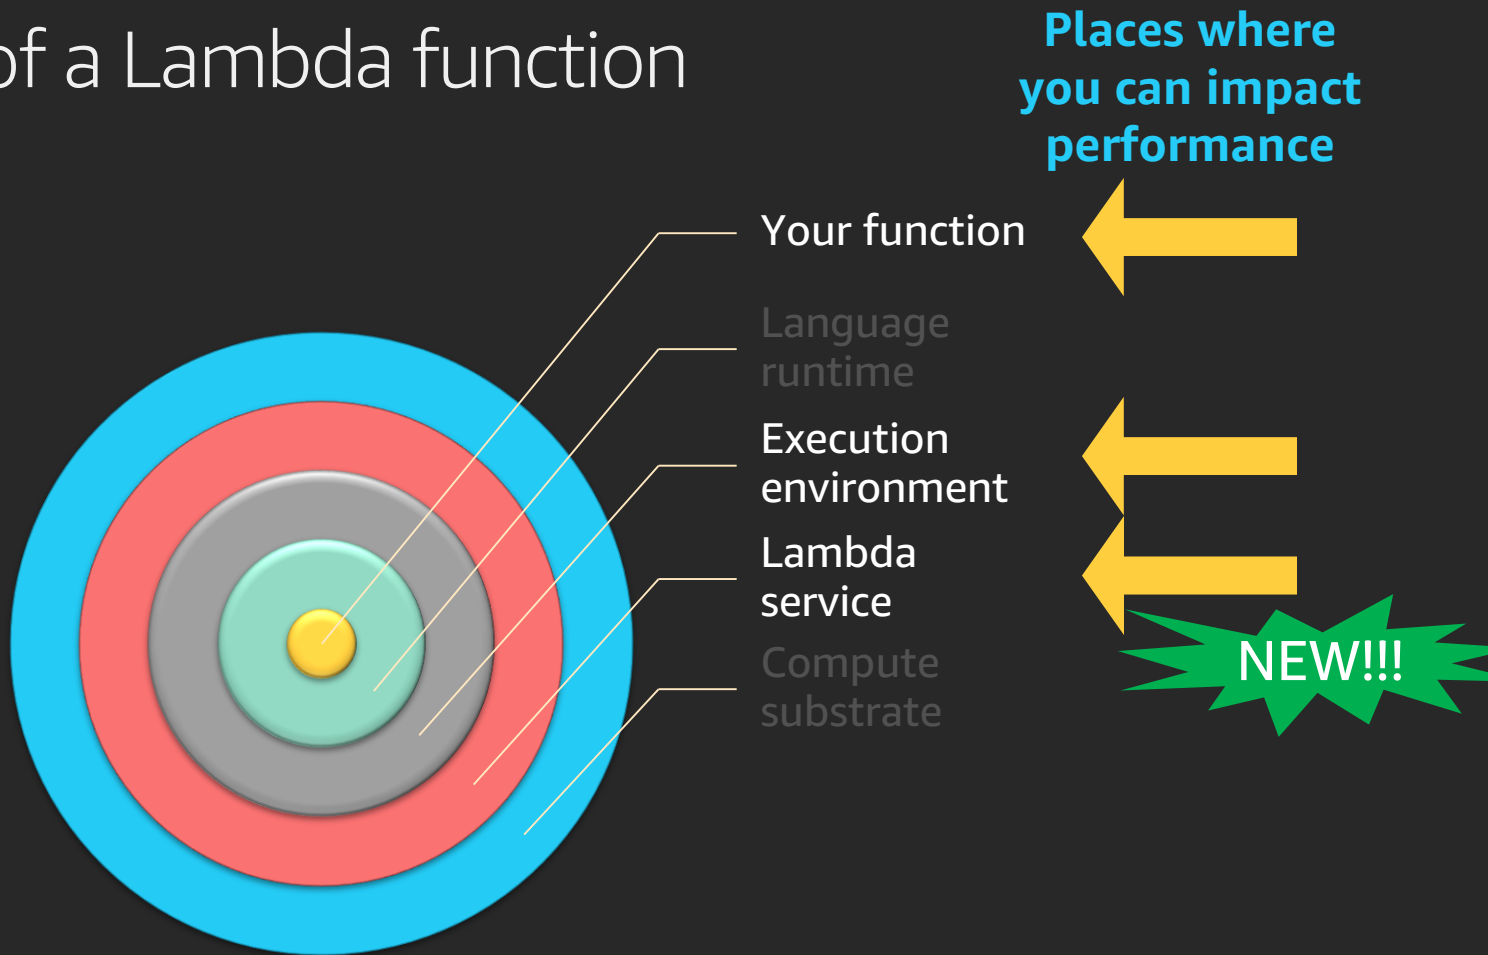


AWS Lambda

# Anatomy of a Lambda function




# Anatomy of a Lambda function





# Anatomy of a Lambda function



# Serverless applications



AWS Lambda

# Serverless applications

## Function



Node.js  
Python  
Java  
C#  
Go  
Ruby  
Runtime API

# Anatomy of a Lambda function

## Handler() function

Function to be executed upon invocation

## Event object

Data sent during Lambda function Invocation

## Context object

Methods available to interact with runtime information (request ID, log group, more)

```
import json

def lambda_handler(event, context):
    # TODO implement
    return {
        'statusCode': 200,
        'body': json.dumps('Hello world!')
    }
```

# Serverless applications

Event source



Function



Changes in  
data state



Requests to  
endpoints



Changes in  
Resource state



Node.js  
Python  
Java  
C#  
Go  
Ruby  
Runtime API

# Anatomy of a Lambda function

```
Function myhandler(event, context) {  
  <Event handling logic> {  
    result = SubfunctionA()  
  }else {  
    result = SubfunctionB()  
  }  
  
  return result;  
}
```

**Your handler**

```
Import sdk
Import http-lib
Import ham-sandwich
```

```
Pre-handler-secret-getter()
Pre-handler-db-connect()
```

```
Function myhandler(event, context) {
  <Event handling logic> {
    result = SubfunctionA()
  }else {
    result = SubfunctionB()

  return result;
}
```

Your handler

```
Import sdk
Import http-lib
Import ...
```

## Dependencies, configuration information, common helper functions

```
Pre-handler-secret-getter()
Pre-handler-db-connect()
```

```
Function myhandler(event, context) {
  <Event handling logic> {
    result = SubfunctionA()
  }else {
    result = SubfunctionB()

  }

  return result;
}
```

Your handler



# Pre-handler code, dependencies, variables

- Import only what you need
  - Where possible trim down SDKs and other libraries to the specific bits required
- Pre-handler code is great for establishing connections, but be prepared to then handle reconnections in further executions
- **REMEMBER** – execution environments are reused
  - Lazily load variables in the global scope
  - Don't load it if you don't need it – cold starts are affected
  - Clear out used variables so you don't run into left-over state

```
Import sdk
Import http-lib
Import ham-sandwich

Pre-handler-secret-getter()
Pre-handler-db-connect()

Function myhandler(event,
context) {
    . . . .
```

```
Import sdk
Import http-lib
Import ...
```

## Dependencies, configuration information, common helper functions

```
Pre-handler-secret-getter()
Pre-handler-db-connect()
```

```
Function myhandler(event, context) {
  <Event handling logic> {
    result = SubfunctionA()
  }else {
    result = SubfunctionB()
  }

  return result;
}
```

## Your handler

```
Function Pre-handler-secret-getter() {
}
```

```
Function Pre-handler-db-connect(){
}
```

```
Import sdk
Import http-lib
Import ...
```

## Dependencies, configuration information, common helper functions

```
Pre-handler-secret-getter()
Pre-handler-db-connect()
```

```
Function myhandler(event, context) {
  <Event handling logic> {
    result = SubfunctionA()
  }else {
    result = SubfunctionB()
  }

  return result;
}
```

### Your handler

```
Function Pre-handler-secret-getter() {
}
```

### Common helper functions

```
Function Pre-handler-db-connect(){
}
```

# AWS Lambda Environment Variables

- Key-value pairs that you can dynamically pass to your function
- Available via standard environment variable APIs such as `process.env` for Node.js or `os.environ` for Python
- Can optionally be encrypted via AWS Key Management Service (AWS KMS)
  - Allows you to specify in IAM what roles have access to the keys to decrypt the information
- Useful for creating environments per stage (i.e., dev, testing, production)



# AWS Systems Manager – Parameter Store

## Centralized store to manage your configuration data

- Supports hierarchies
- Plaintext or encrypted with AWS KMS
- Can send notifications of changes to Amazon SNS/AWS Lambda
- Can be secured with IAM
- Calls recorded in AWS CloudTrail
- Can be tagged
- Works with AWS Secrets Manager
- Available via API/SDK

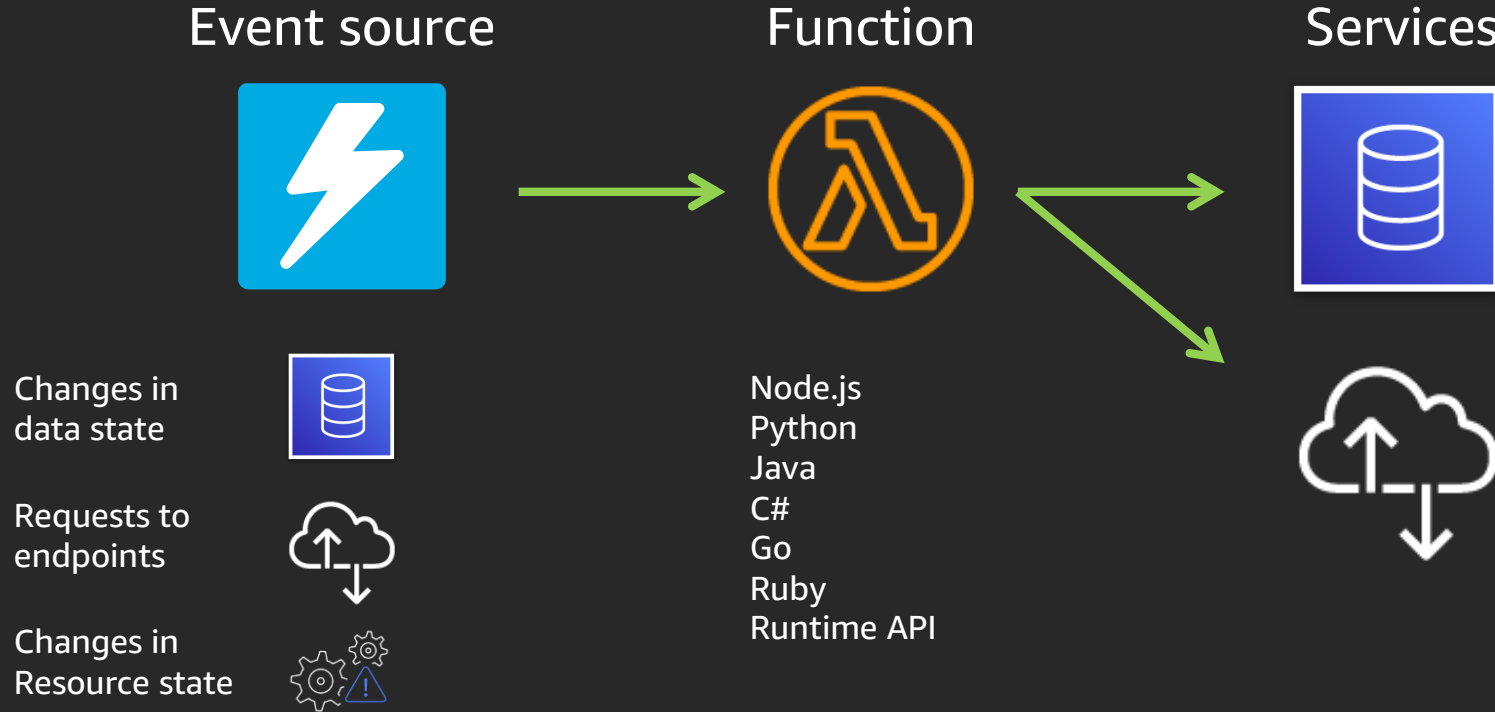
Useful for centralized environment variables, secrets control, feature flags

```
from __future__ import print_function
import json
import boto3
ssm = boto3.client('ssm', 'us-east-1')

def get_parameters():
    response = ssm.get_parameters(
        Names=['LambdaSecureString'], withDecryption=True
    )
    for parameter in response['Parameters']:
        return parameter['Value']

def lambda_handler(event, context):
    value = get_parameters()
    print("value1 = " + value)
    return value # Echo back the first key value
```

# Serverless applications



```
Import sdk
Import http-lib
Import ...
```

## Dependencies, configuration information, common helper functions

```
Pre-handler-secret-getter()
Pre-handler-db-connect()
```

```
Function myhandler(event, context) {
  <Event handling logic> {
    result = SubfunctionA()
  }else {
    result = SubfunctionB()

  }

  return result;
}
```

### Your handler

```
Function Pre-handler-secret-getter() {
}
```

### Common helper functions

```
Function Pre-handler-db-connect(){
}
```

```
Function subFunctionA(thing){
  ## logic here
}
```

```
Function subFunctionB(thing){
  ## logic here
}
```

```
Import sdk
Import http-lib
Import ...
```

## Dependencies, configuration information, common helper functions

```
Pre-handler-secret-getter()
Pre-handler-db-connect()
```

```
Function myhandler(event, context) {
  <Event handling logic> {
    result = SubfunctionA()
  }else {
    result = SubfunctionB()
  }

  return result;
}
```

## Your handler

```
Function Pre-handler-secret-getter() {
}
```

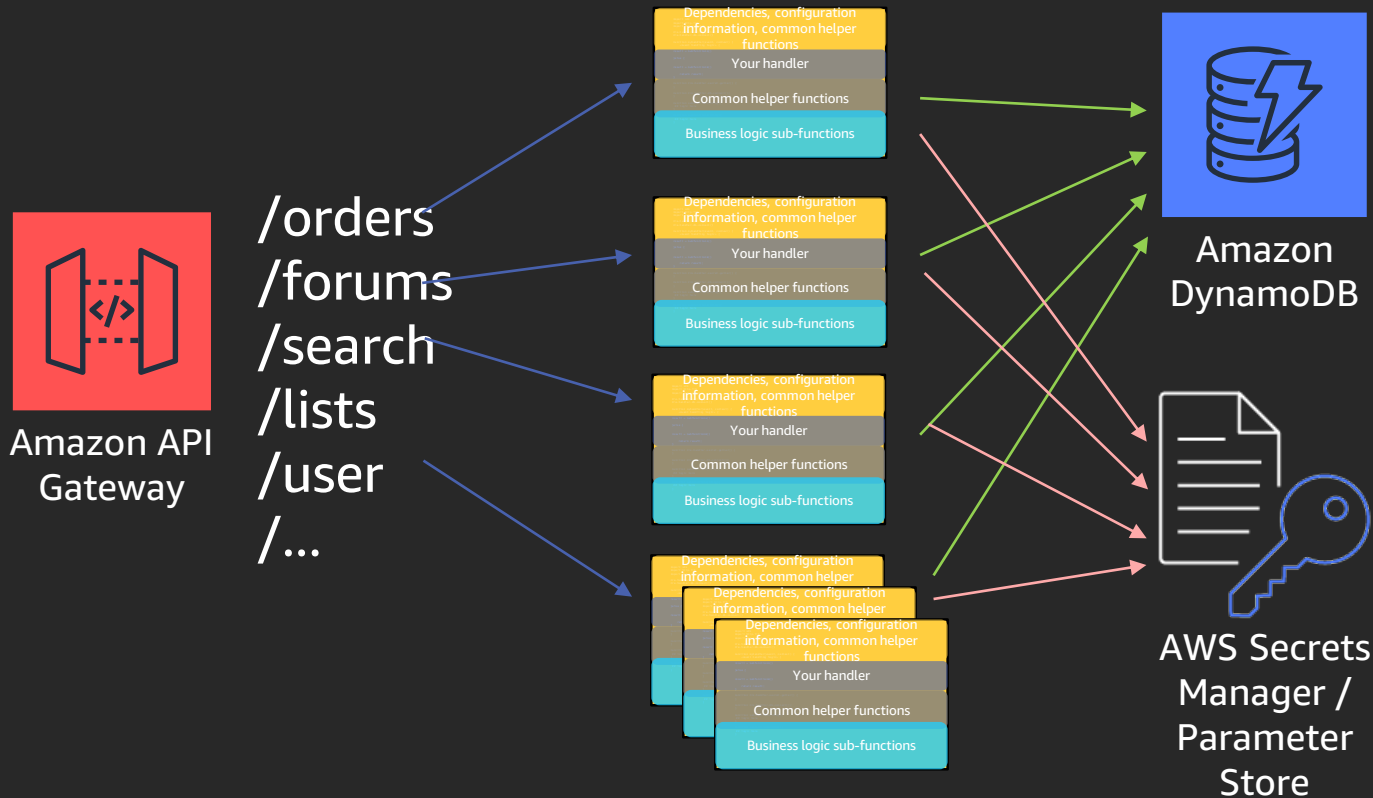
## Common helper functions

```
Function Pre-handler-db-connect(){
}
```


## Business logic subfunctions



# Anatomy of a serverless application








# Anatomy of a serverless application




Amazon API  
Gateway

/orders  
/forums  
/search  
/lists  
/user  
/...

There could be a lot  
of duplicated code  
here!



Amazon  
DynamoDB



AWS Secrets  
Manager /  
Parameter  
Store

# Lambda Layers



Lets functions easily share code: Upload layer once, reference within any function

Layer can be anything: dependencies, training data, configuration files, etc

Promote separation of responsibilities, lets developers iterate faster on writing business logic

Built-in support for secure sharing by ecosystem

# Using Lambda Layers

- Put common components in a ZIP file and upload it as a Lambda layer
- Layers are immutable and can be versioned to manage updates
- When a version is deleted or permissions to use it are revoked, functions that used it previously will continue to work, but you won't be able to create new ones
- You can reference up to five layers, one of which can optionally be a custom runtime



Lambda  
Layers

arn:aws:lambda:region:accountId:layer:shared-lib:1



Lambda  
Layers

arn:aws:lambda:region:accountId:layer:shared-lib:2



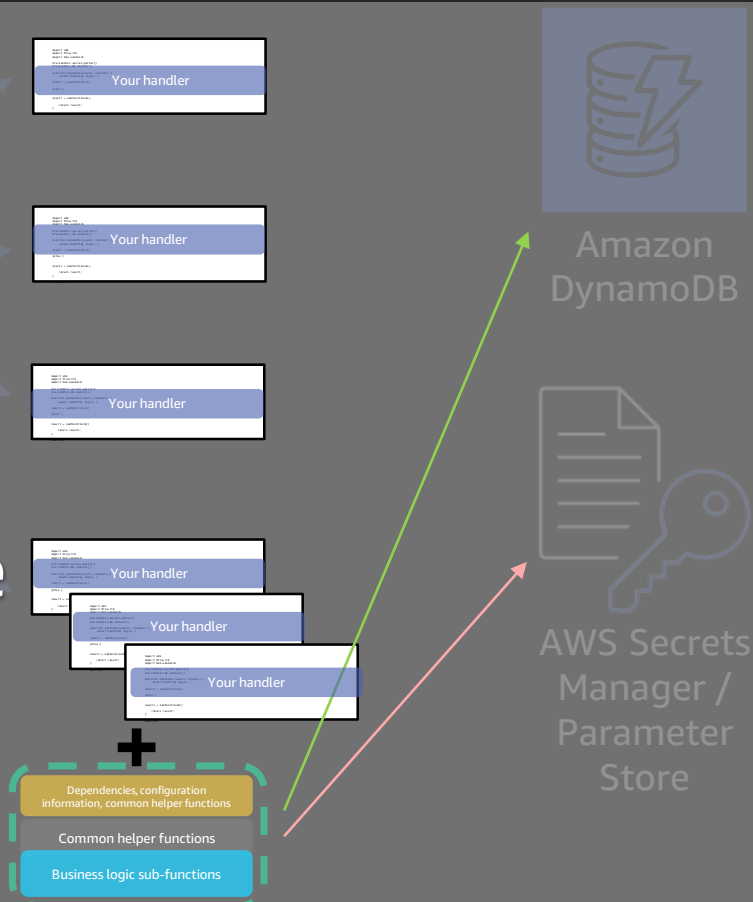
Lambda  
Layers

arn:aws:lambda:region:accountId:layer:shared-lib:3

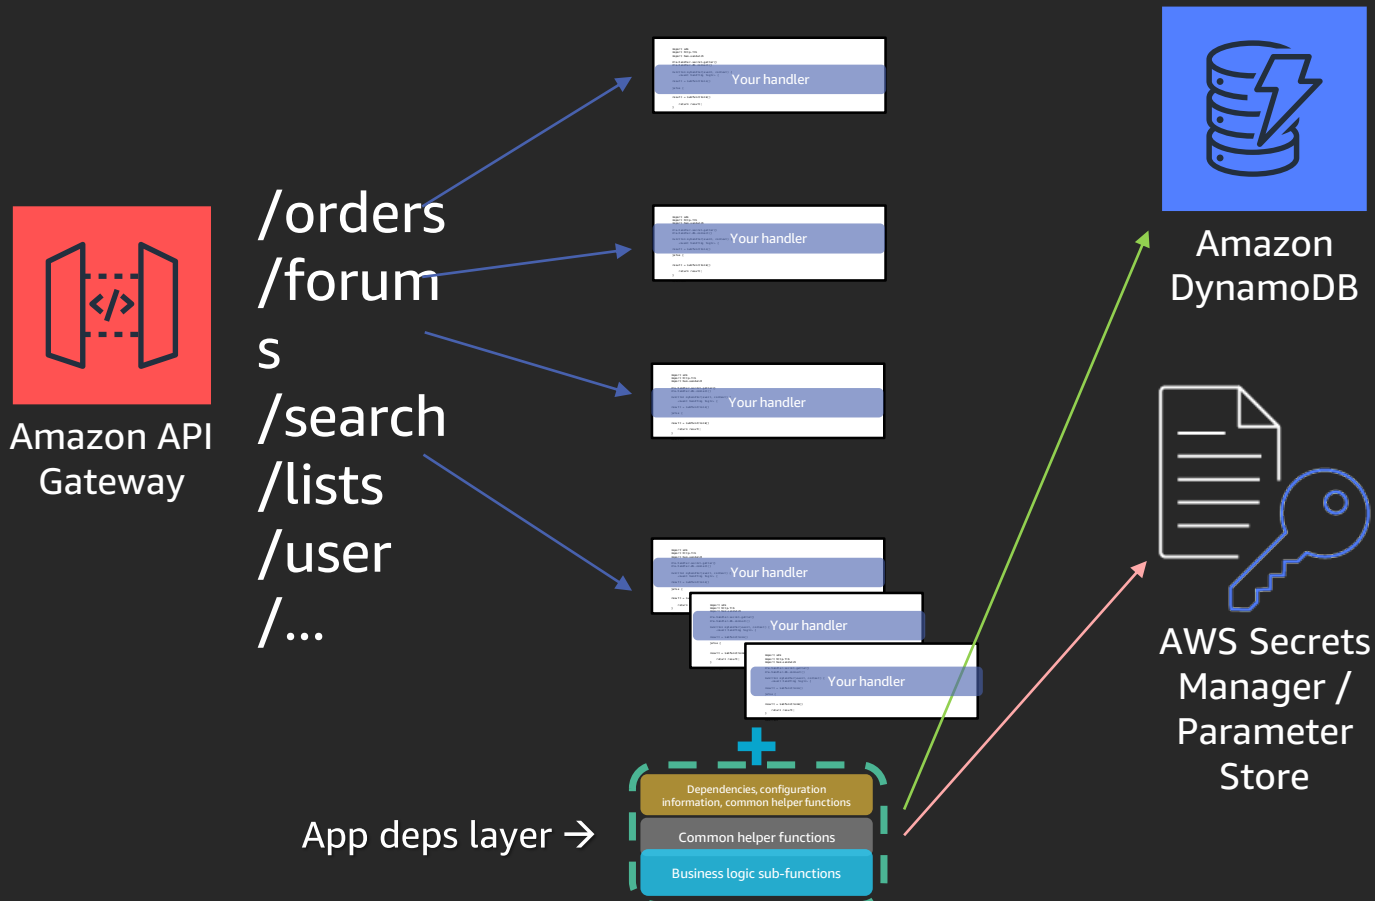
# Anatomy of a serverless application

With Lambda Layers you can reduce the duplication of code


App deps layer →



# Anatomy of a serverless application



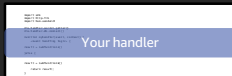
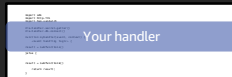
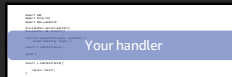


# Anatomy of a serverless application



Amazon API  
Gateway

/orders  
/forum  
s  
/search  
/lists  
/user  
/...



App deps layer →

What if we wanted  
to use a different  
database?

Like Amazon RDS?

DynamoDB



AWS Secrets  
Manager /  
Parameter  
Store

# Introducing Amazon RDS Proxy

NEW!!!

## Simplifies connecting to Amazon RDS databases from Lambda

- Reduces connections via a shared connection pool to your database
- Integrates with Secrets Manager for simple authentication
- Handles failover of database instances transparently for you
- In preview today: limited regions, supports just MySQL

### Database proxies (preview)

Add database proxy

Proxy identifier

Status

Engine compatibility

prod-db-proxy [🔗](#)

🟢 Available

MYSQL

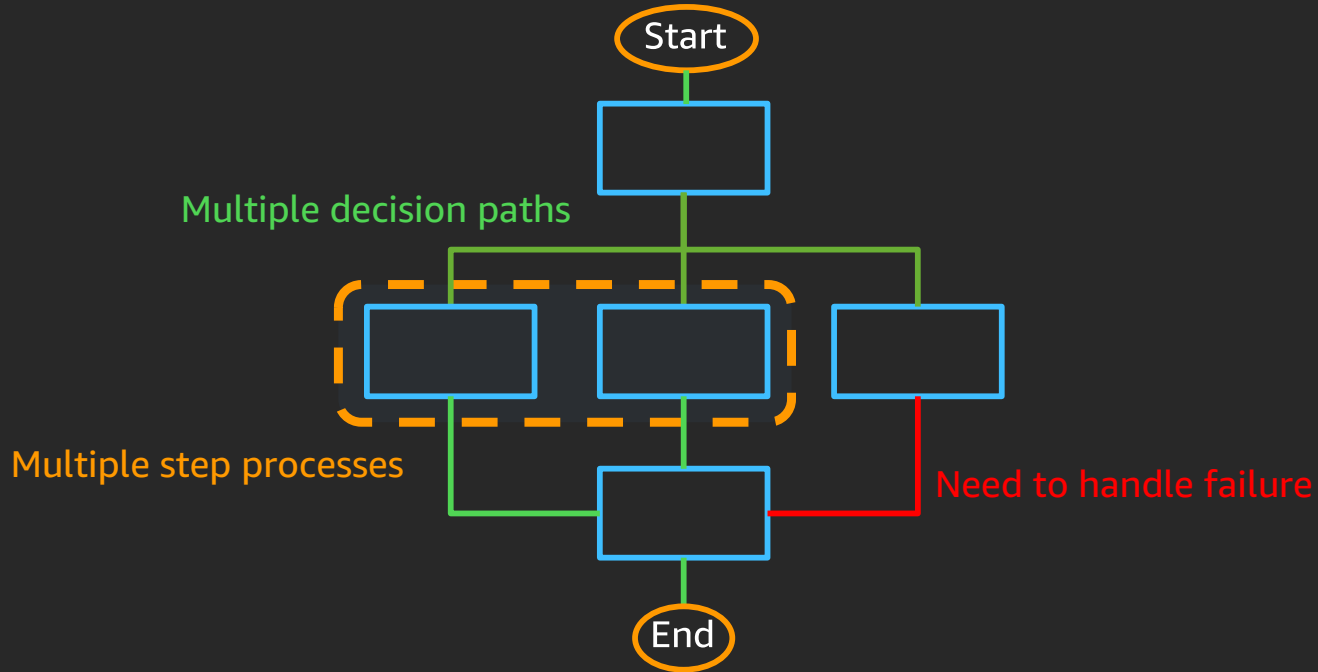


# Less code > more code

# Concise function logic

- Use functions to **TRANSFORM**, not **TRANSPORT**
  - Use purposeful built services for communication fan-out, message handling, data replication, writing to data stores/databases
- Leave retry and error handling to the services themselves
- Read only what you need. For example:
  - Message filters in Amazon SNS
  - Fine grained rules in Amazon EventBridge
  - Query filters in Amazon RDS Aurora
  - Use Amazon S3 Select
  - Properly indexed databases

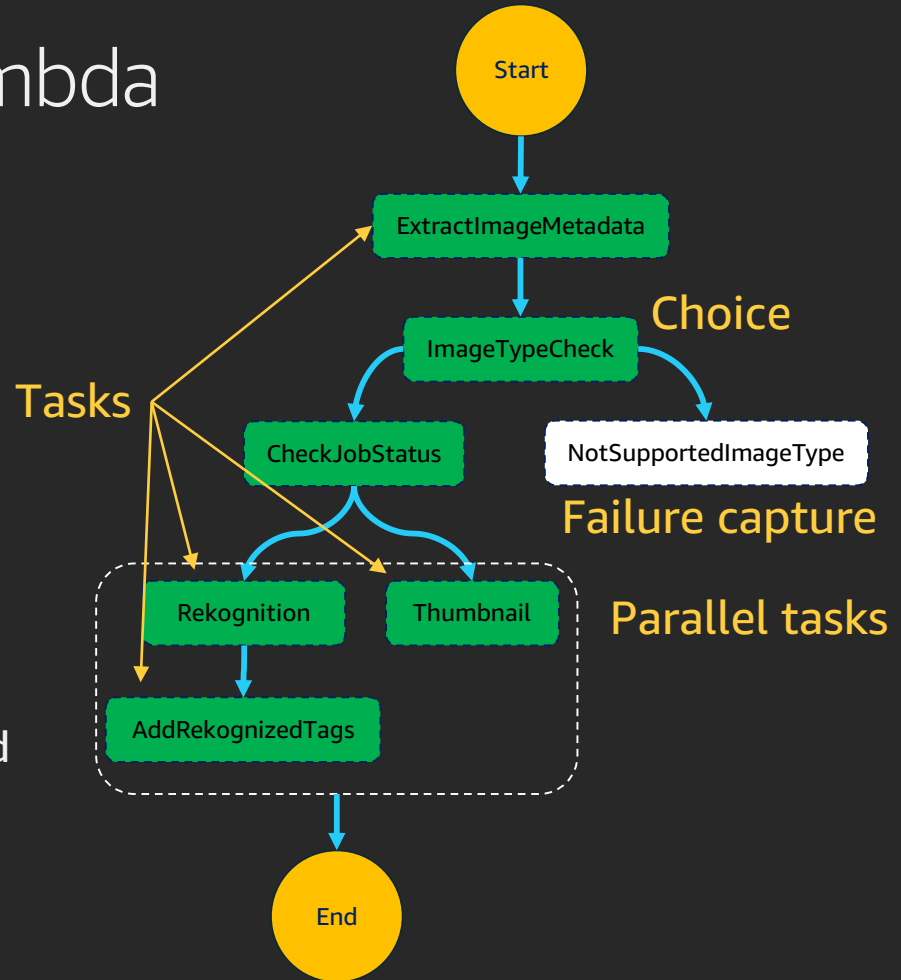
# Business workflow is rarely sequential start to finish



# AWS Step Functions + Lambda

## “Serverless” workflow management with zero administration:

- Makes it easy to coordinate the components of distributed applications and microservices using visual workflows
- Automatically triggers and tracks each step and retries when there are errors, so your application executes in order and as expected
- Logs the state of each step, so when things do go wrong, you can diagnose and debug problems quickly





# Step Functions: Integrations



Simplify building workloads such as order processing, report generation, and data analysis

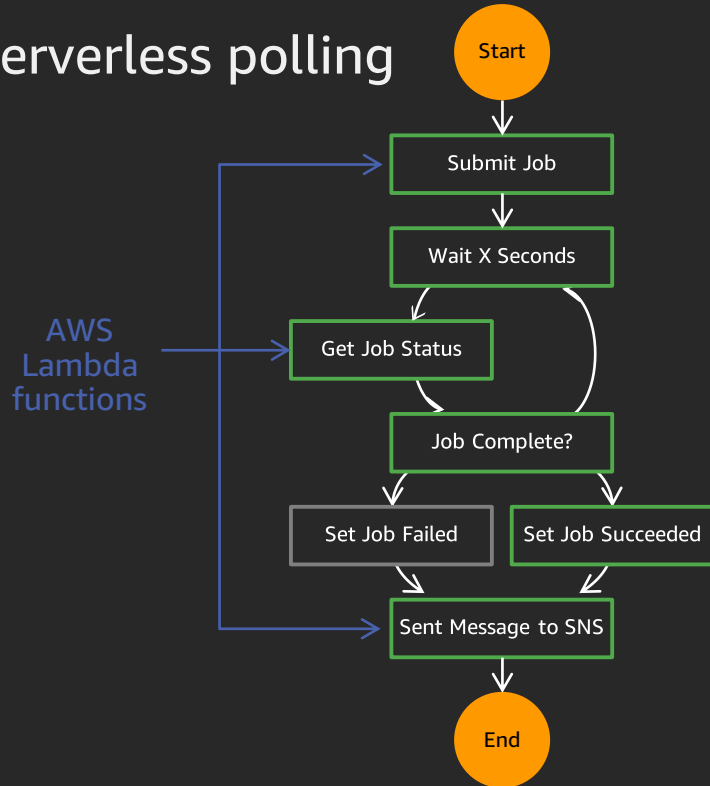
Write and maintain less code; add services in minutes

More service integrations:

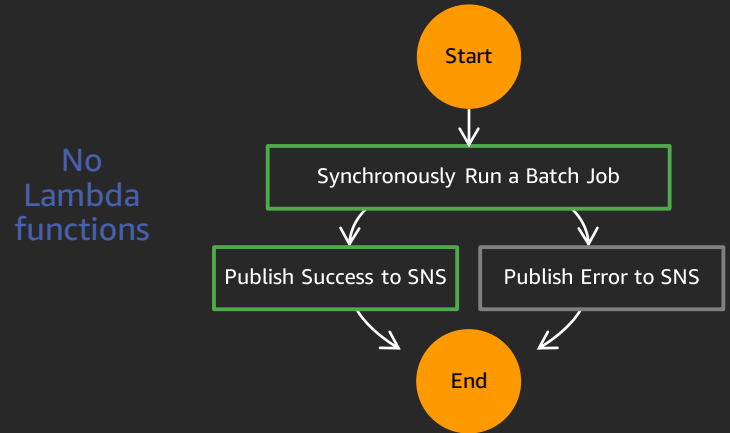


# Simpler integration, less code

## With serverless polling



## With direct service integration



# Introducing Lambda Event Destinations

NEW!!!

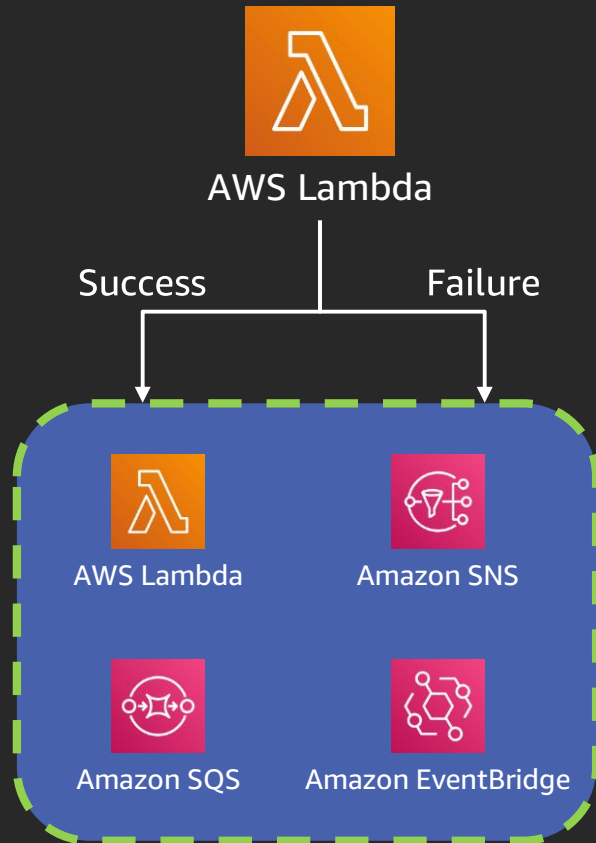
For asynchronous invocations, capture success or failure

- Record contains details about the request and response in JSON format
- Contains more information than data sent to a DLQ

- Can send both outcomes to same destination

or


- Can send success to one destination, failure to another



**The best performing Lambda function is the one you rip out and replace with a built in integration**




# Anatomy of a Lambda function



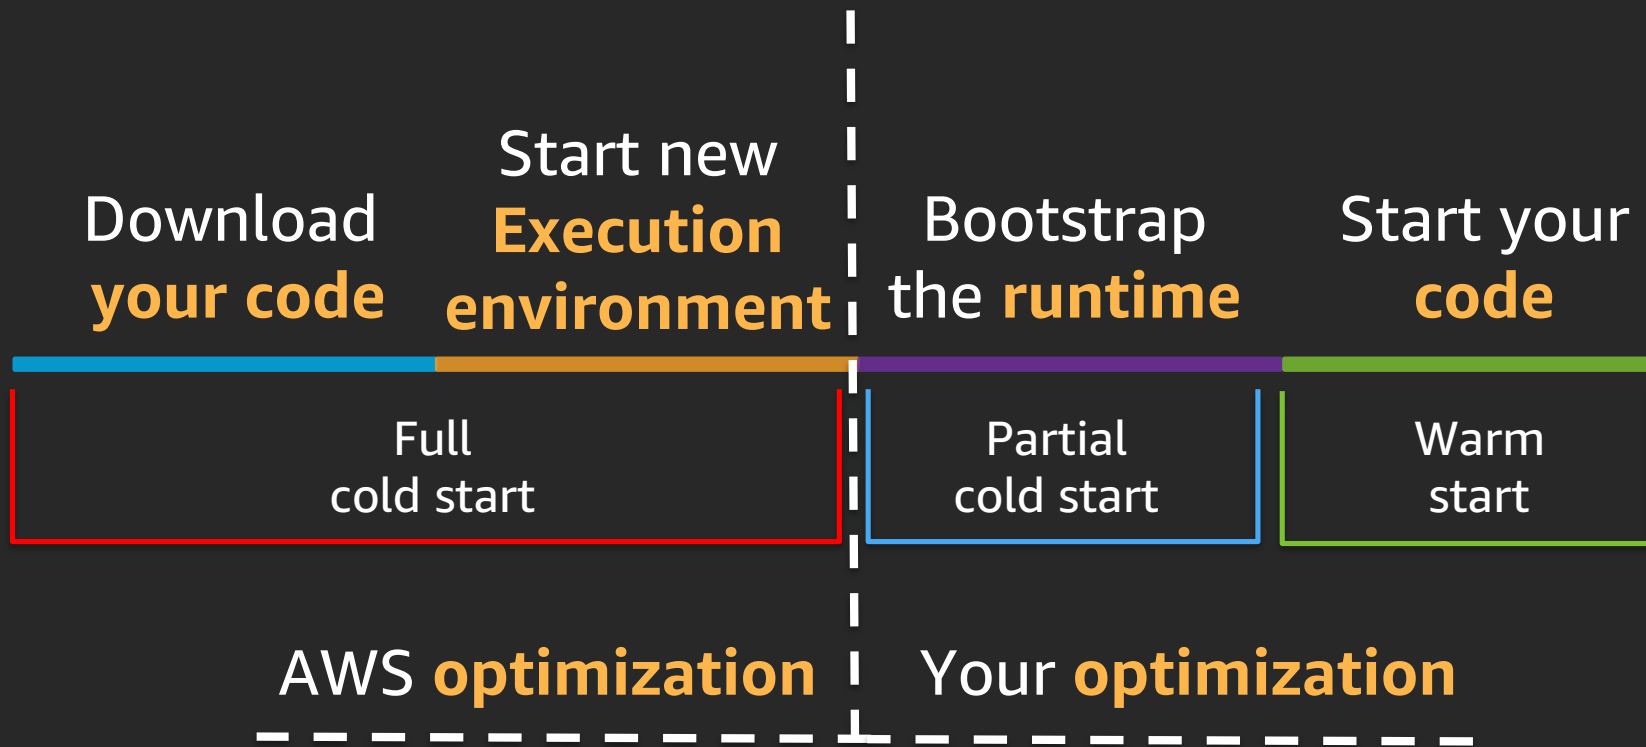
## Recap:

- Minimize dependencies
- Use pre-handler logic sparingly but strategically
- Share secrets based on application scope:
  - Single function: Env-Vars
  - Multi Function/shared environment: Parameter Store
- Think about how re-use impacts variables, connections, and dependency usage
- Layers save on code duplication and help enable standardization across functions
- Amazon RDS Proxy will simplify relational database usage with Lambda
- Concise logic
- Push orchestration up to Step Functions
- Lambda destinations will simplify and improve asynchronous workflows

# Anatomy of a Lambda function



# The function lifecycle

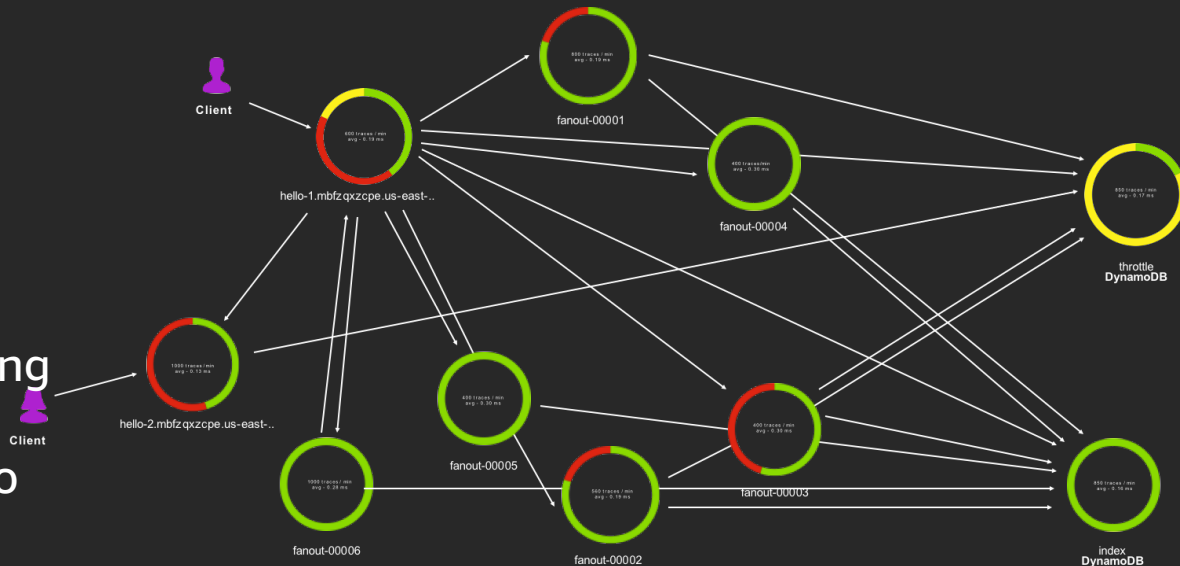


# AWS X-Ray

Profile and troubleshoot serverless applications:

- Lambda instruments incoming requests for all supported languages and can capture calls made in code
- API Gateway inserts a tracing header into HTTP calls as well as reports data back to X-Ray itself

```
var AWSXRay = require('aws-xray-sdk-core');  
var AWS = AWSXRay.captureAWS(require('aws-sdk'));  
S3Client = AWS.S3();
```



Enable X-Ray Tracing ☒

Enable active tracing [Info](#)



# X-Ray Trace Example


Method	Response	Duration	Age	ID
--	202	2.0 sec	1.3 min (2017-04-14 00:42:54 UTC)	1-58f01b0e-53eef2bd463eecd7f311ce4

Name	Res.	Duration	Status	0.0ms	200ms	400ms	600ms	800ms	1.0s	1.2s	1.4s	1.6s	1.8s	2.0s
------	------	----------	--------	-------	-------	-------	-------	-------	------	------	------	------	------	------


▼ **s3example** AWS::Lambda

Operation	Count	Time	Status	Visual
s3example	202	87.0 ms	✓	
Dwell Time	-	186 ms	✓	
Attempt #1	200	1.8 sec	✓	

▼ **s3example** AWS::Lambda::Function

Component	Count	Time (ms)	Status	Visual
s3example	-	863 ms	✓	
Initialization	-	334 ms	✓	
S3	404	762 ms	!	 PutObject

# Tweak your function's computer power



Lambda exposes only a memory control, with the **% of CPU core and network capacity** allocated to a function proportionally

Is your code CPU, Network, or memory-bound? If so, it could be **cheaper** to choose more memory.

# Smart resource allocation

Match resource allocation (up to 3 GB!) to logic

Stats for Lambda function that calculates **1000 times** all prime numbers **<= 1000000**

<b>128 MB</b>	<b>11.722965sec</b>	<b>\$0.024628</b>
<b>256 MB</b>	<b>6.678945sec</b>	<b>\$0.028035</b>
<b>512 MB</b>	<b>3.194954sec</b>	<b>\$0.026830</b>
<b>1024 MB</b>	<b>1.465984sec</b>	<b>\$0.024638</b>

**Green**==Best

**Red**==Worst

# Smart resource allocation

Match resource allocation (up to **3 GB!**) to logic

Stats for Lambda function that calculates **1000 times** all prime numbers **<= 1000000**

**128 MB**

**256 MB**

**512 MB**

**1024 MB**

11.722965sec  
6.678945sec  
**-10.256981sec**  
3.194954sec  
1.465984sec

\$0.024628  
\$0.028035  
**+\$0.00001**  
\$0.026830  
\$0.024638

**Green**==Best

**Red**==Worst



# Multithreading? Maybe!

- <1.8 GB is still single core

- CPU bound workloads won't see gains – processes share same resources

- >1.8 GB is multicore

- CPU bound workloads will gain, but need to multithread

- I/O bound workloads WILL likely see gains

- e.g. parallel calculations to return

**AWS Compute Blog**

by [Chris Munns](#) | on 03 SEP 2019 | in [Amazon VPC](#), [AWS Lambda](#), [Serverless](#) | [Permalink](#) | [Comments](#) | [Share](#)

**NEW!!!**

← Before: 14.8-sec duration

After: 933 ms duration →

Timeline

Raw data

Method	Response	Duration	Age	ID
--	200	933 ms	52.1 sec (2019-08-06 18:12:12 UTC)	1-5d49c2fc-82899913a8dc997e71c6352a

Name	Res.	Duration	Status	0.0ms	100ms	200ms	300ms	400ms	500ms	600ms	700ms	800ms	900ms	1.0s
▼ internet-access AWS::Lambda														
internet-access	200	933 ms	✓											
▼ internet-access AWS::Lambda:Function														
internet-access	-	495 ms	✓											
Initialization	-	167 ms	✓											
Invocation	-	456 ms	✓											
Overhead	-	39.2 ms	✓											

# Lambda + VPC, no longer a cold-start pain point!

AWS Compute Blog

## Announcing improved VPC networking for AWS Lambda functions

by Chris Munns | on 03 SEP 2019 | in Amazon VPC, AWS Lambda, Serverless | Permalink | Comments | Share

NEW!!!

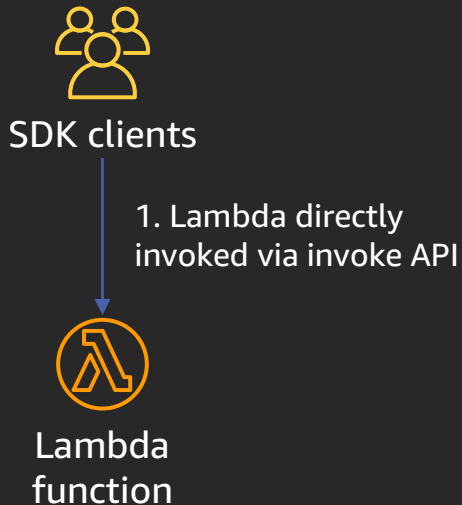
# Now available in \*all\* public regions

After: 933ms duration →

Method	Response	Duration	Age	ID
--	200	933 ms	52.1 sec (2019-08-06 18:00:00 UTC)	1-5d49c2fc-52899913a8dc997e71c6352a
Name				
Res.				
Duration				
Status				
0.0ms100ms200ms300ms400ms500ms600ms700ms800ms900ms100ms				
internet-access	AWS Lambda			
internet-access	200	933 ms	✓	
internet-access	AWS Lambda Function			
internet-access	-	495 ms	✓	
Initialization	-	167 ms	✓	
Invocation	-	456 ms	✓	
Overhead	-	39.2 ms	✓	

# Events and you

# Lambda API



API provided by the Lambda service

Used by all other services that invoke Lambda across all models


Supports sync and async

Can pass any event payload structure you want


Client included in every SDK

# Lambda execution model

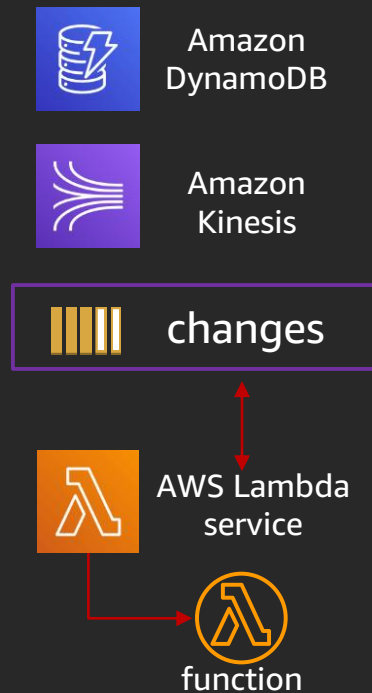
## Synchronous (push)




## Asynchronous (event)




## Stream (Poll-based)




# Synchronous APIs



# Synchronous APIs




retry after  
failure






# Synchronous APIs



# Synchronous APIs



# Synchronous APIs




The diagram illustrates a synchronous API interaction. On the left, a 'Client' icon is shown. A blue arrow labeled 'POST /order' points from the Client to a green box labeled 'Order'. From the 'Order' box, another blue arrow labeled 'POST /invoice' points to a blue box labeled 'Invoice'. Below these arrows, green arrows point back from the 'Order' and 'Invoice' boxes to the 'Client'. Two blue starburst shapes, representing errors or failures, are placed on the return green arrows. The text 'Who owns the retry? For how long? Does the client ever know? Etc..' is overlaid on the diagram.

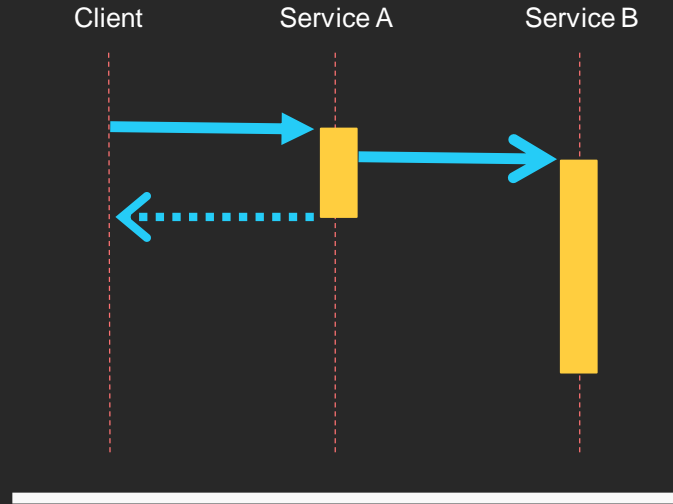
Who owns the retry? For how long?  
Does the client ever know? Etc..

This effectively creates a “tight coupling”  
where failures become harder to recover from

# Thinking asynchronously

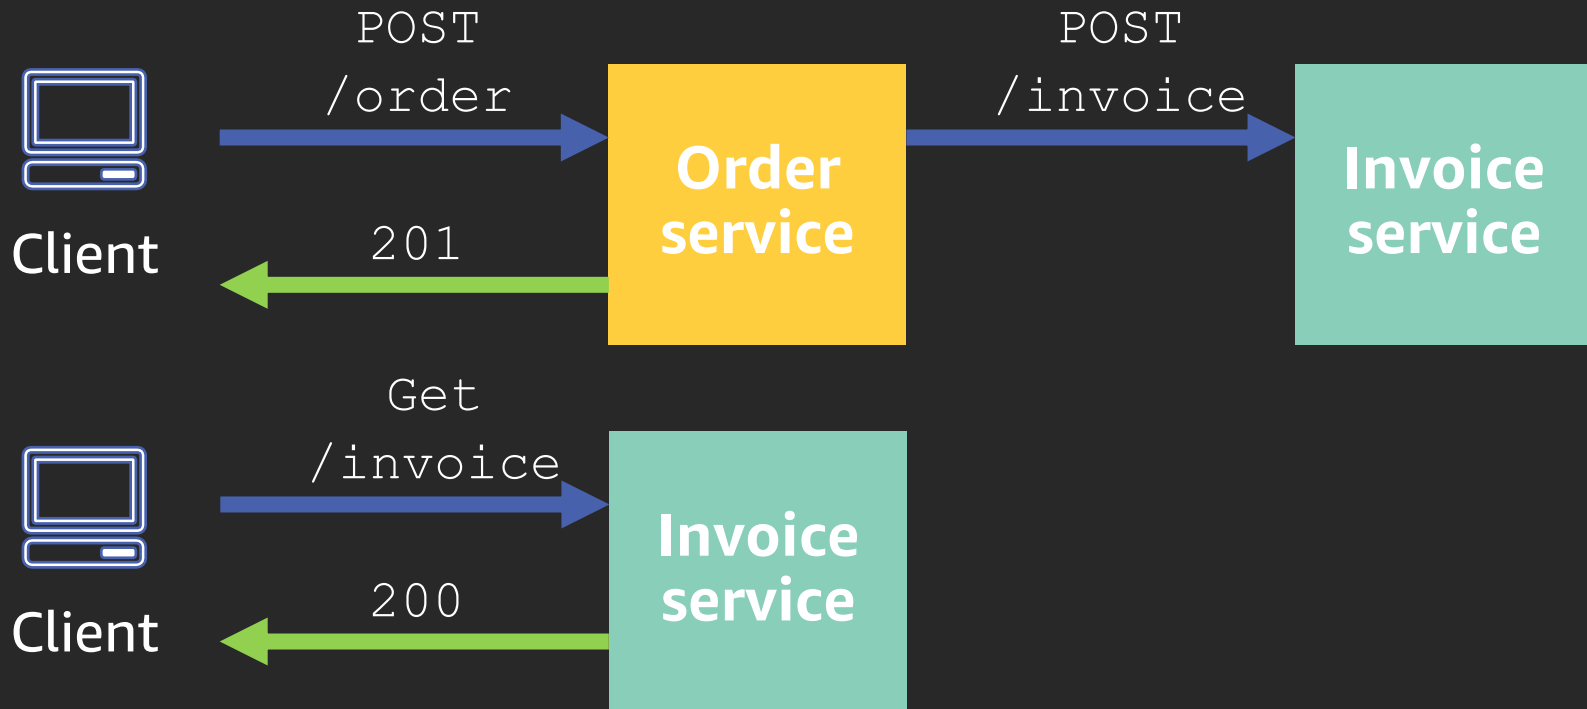


Synchronous  
commands



Asynchronous  
events

# Asynchronous APIs



If you don't need a response, execute async

## Use the Lambda APIs to start an asynchronous execution


- Built-in queue (Amazon SQS behind the scenes)
- Automatic retries
- Dead letter queue for failed events

```
client =  
boto3.client("lambda")  
  
client.invoke_async(  
    FunctionName="test"  
    InvokeArgs=json_payload  
)
```


# Topics, streams, queues, and buses



Amazon SNS



Amazon SQS



Amazon EventBridge




Amazon Kinesis  
Data Streams

# Ways to compare




Scale/concurrency  
controls



Durability



Persistence



Consumption  
models




Retries



Pricing



# Concurrency across models




# Recent announcements for async event sources

- Amazon SQS FIFO as invoke source for Lambda
- Amazon SNS Dead Letter Queues (DLQs)
- For streamed events:
  - MaximumRetryAttempts, MaximumRecordAgeInSeconds, BisectBatchOnFunctionError, On-failure destination
  - Batch Window
  - Parallelization Factor
- For async events:
  - MaximumRetryAttempts
  - MaximumEventAgeInSeconds




ICYMI: Serverless  
pre:Invent 2019

# Directed vs. Observable events



Directed  
Commands

# Directed vs. Observable events

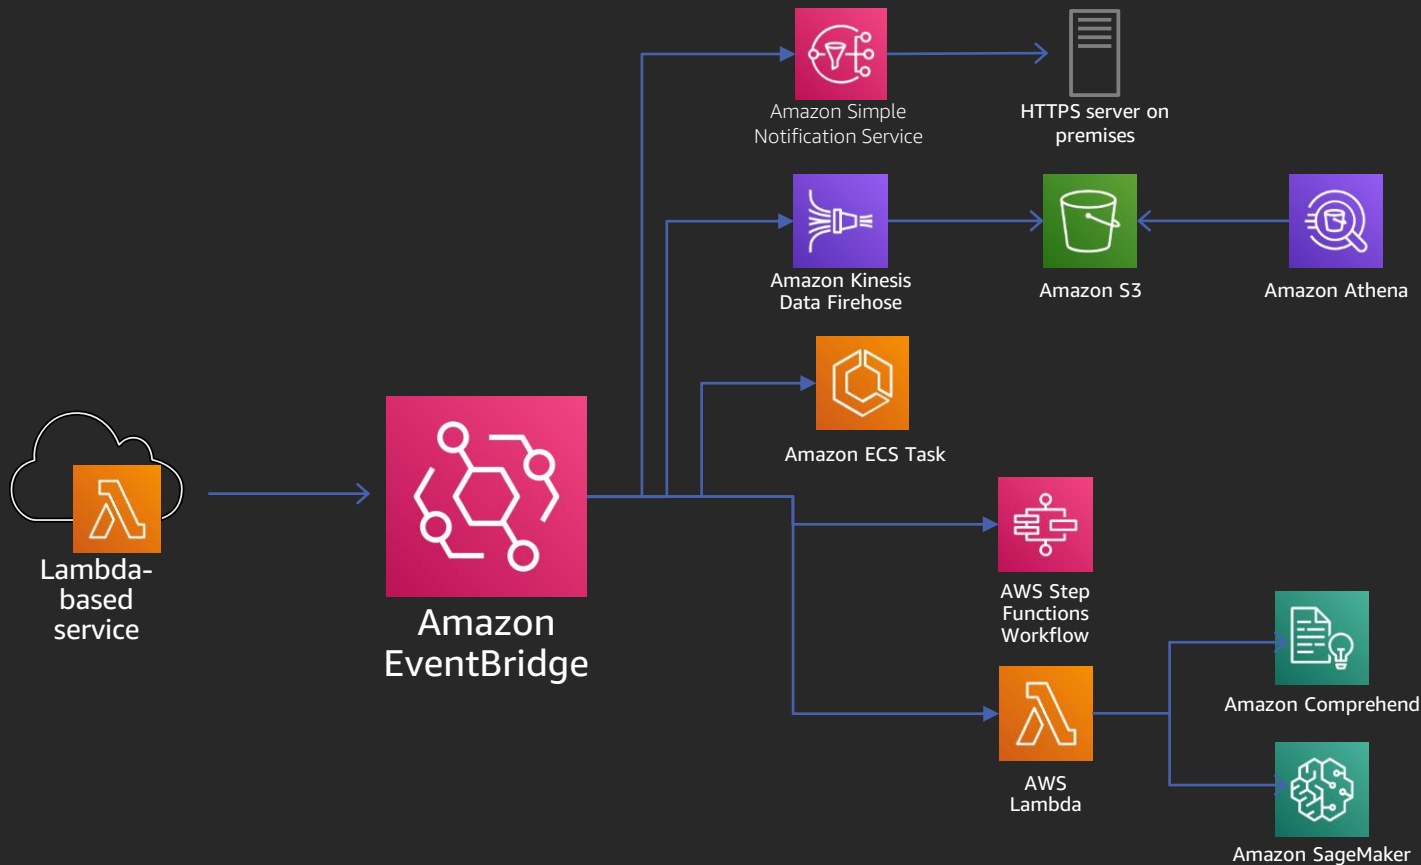


Directed  
Commands




Observable  
Events

# Event passing with EventBridge



# Events with EventBridge



- Your services can both produce messages onto the bus and consume just the messages they need from the bus
- Services don't need to know about each other, just about the bus.


# Recent announcements for async event sources

AWS Compute Blog

## Introducing Amazon EventBridge schema registry and discovery – In preview


by Julian Wood | on 01 DEC 2019 | in Amazon EventBridge, Serverless | Permalink | Comments | Share

NEW!!!



# Lambda per function concurrency controls

- Concurrency a shared pool by default
- Separate using per function concurrency settings
  - Acts as reservation
- Also acts as max concurrency per function
  - Especially critical for downstream resources like databases
- “Kill switch” – set per function concurrency to zero





# Lambda Dead Letter Queues

“By default, a failed Lambda function **invoked asynchronously is retried twice**, and then the event is discarded.” –

<https://docs.aws.amazon.com/lambda/latest/dg/dlq.html>

- **Turn this on!** (for async use-cases)
- Monitor it via an **SQS Queue length metric/alarm**
- If you use SNS, send the messages to something durable and/or a trusted endpoint for processing
  - Can send to Lambda functions in other regions
- If and when things go “**boom**” DLQ can save your invocation event information



In many ways, Lambda Destinations supersedes this

**Friends don't let friends  
"Action": "s3:\*"**

# Lambda permissions model

- **Function policies:**

- Example: "Actions on bucket X can invoke Lambda function Z"
- Resource policies allow for cross-account access
- Used for sync and async invocations

- **Execution role:**

- Example: "Lambda function A can read from DynamoDB table users"
- Define what AWS resources/API calls this function can access via IAM
- Used in streaming invocations



# AWS Serverless Application Model (AWS SAM)



AWS CloudFormation extension optimized for serverless

Special serverless resource types: functions, APIs, tables, layers, and applications

Supports anything AWS CloudFormation supports

Open specification (Apache 2.0)

<https://aws.amazon.com/serverless/sam>

# AWS SAM template

AWSTemplateFormatVersion: '2010-09-09'

Transform: AWS::Serverless-2016-10-31

Resources:

GetProductsFunction:

Type: AWS::Serverless::Function

Properties:

Handler: index.getProducts

Runtime: nodejs8.10

CodeUri: src/

Policies:

- DynamoDBReadPolicy:

TableName: !Ref ProductTable

Events:

GetResource:

Type: Api

Properties:

Path: /products/{productId}


Method: get

ProductTable:

Type: AWS::Serverless::SimpleTable

Just 20 lines to create:

- Lambda function
- IAM role
- API Gateway
- DynamoDB table



# AWS SAM policy templates

GetProductsFunction:

Type: `AWS::Serverless::Function`

Properties:

...

Policies:

- DynamoDBReadPolicy:

TableName: `!Ref ProductTable`

...

ProductTable:

Type: `AWS::Serverless::SimpleTable`




```
3  "Templates": {
4    "SQSPollerPolicy": {
5      "Description": "Gives permissions to poll an SQS Queue",
6      "Parameters": {
7        "QueueName": {
8          "Description": "Name of the SQS Queue"
9        }
10     },
11     "Definition": {
12       "Statement": [
13         {
14           "Effect": "Allow",
15           "Action": [
16             "sqs:ChangeMessageVisibility",
17             "sqs:ChangeMessageVisibilityBatch",
18             "sqs:DeleteMessage",
19             "sqs:DeleteMessageBatch",
20             "sqs:GetQueueAttributes",
21             "sqs:ReceiveMessage"
22           ],
23           "Resource": {
24             "Fn::Sub": [
25               "arn:${AWS::Partition}:sqs:${AWS::Region}:${AWS::AccountId}:${queueName}",
26               {
27                 "queueName": {
28                   "Ref": "QueueName"
29                 }
30             ]
31           }
32         }
33       ]
34     }
35   }
36 }
```

**50+ predefined  
policies**

**All found here:  
<https://bit.ly/2xWycnj>**

# Anatomy of a Lambda function




## Recap:

- More memory == More CPU and I/O (proportionally)
  - Can also be lower cost
- Use AWS X-Ray to profile your workload
- >1.8GB memory get's you 2 cores, but you might not use/need it
- Think deeply about your execution model and invocation source needs
  - Not everything needs to be an API
- Thinking async will get you over some of the biggest scaling challenges
- Understand the various aspects to queues, topics, streams, and event buses when using them
- Minimize the scope of IAM permissions
  - Leverage tooling like AWS SAM



# Anatomy of a Lambda function




# Introducing Provisioned Concurrency for Lambda

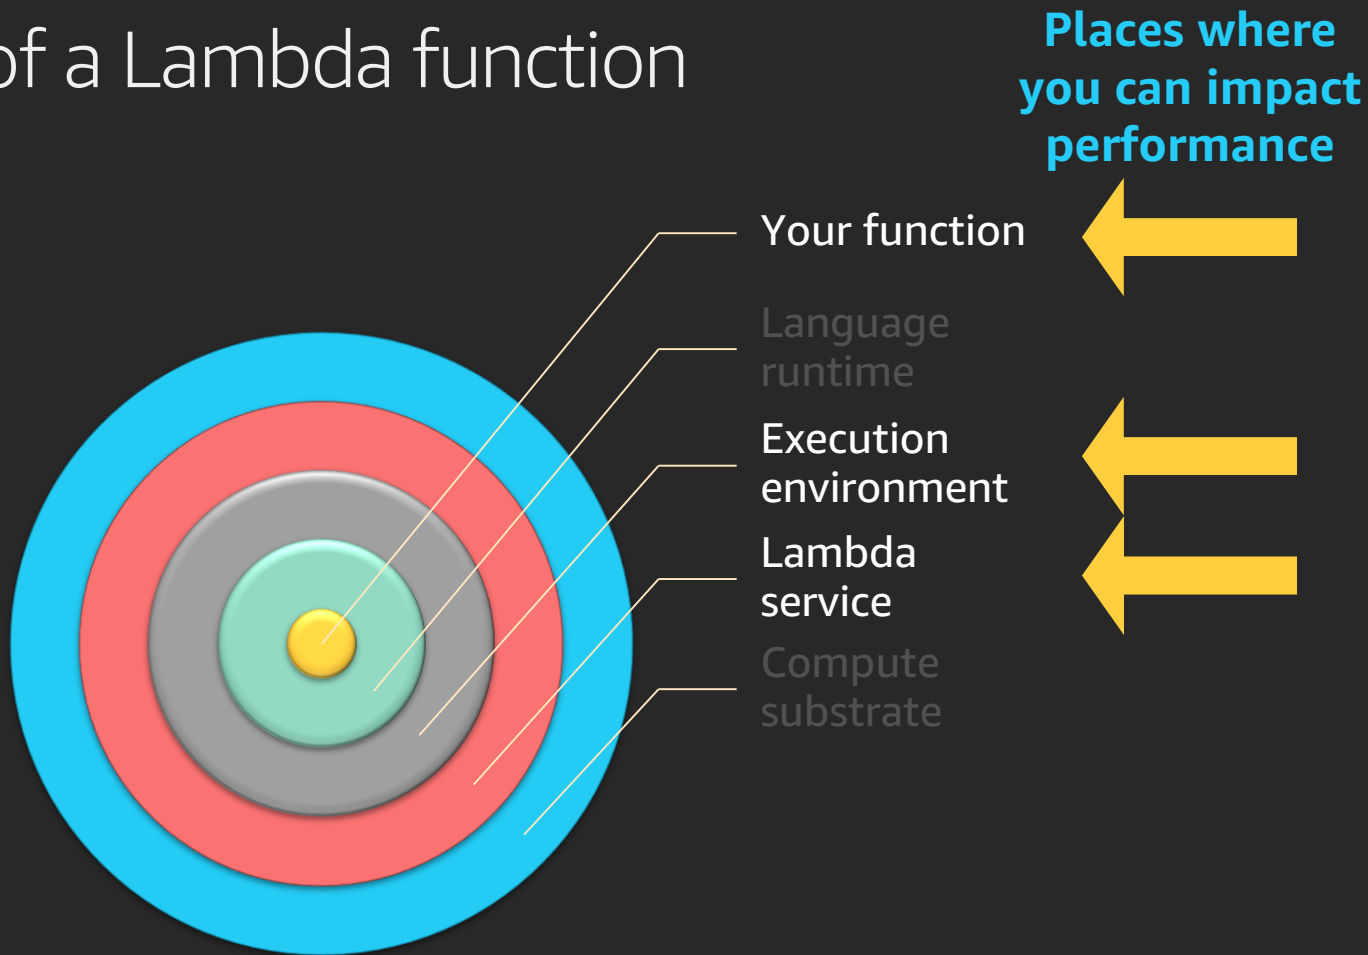


Pre-creates execution environments all the way up through the INIT phase.

- Mostly for interactive workloads that are heavily latency sensitive
- Greatly improved consistency across the full long tail of performance
- Little to no changes to your code or way you use Lambda
- Works with AWS Auto Scaling
- Adds a cost factor for per concurrency provisioned but a lower duration cost for execution
  - This could end up saving you money when heavily utilized



# Anatomy of a Lambda function



# FIN/ACK

## Your Function Recap:

- Minimize dependencies
- Use pre-handler logic sparingly but strategically
- Share secrets based on application scope:
  - Single function: Env-Vars
  - Multi Function/shared environment: Parameter Store
- Think about how re-use impacts variables, connections, and dependency usage
- Layers save on code duplication and help enable standardization across functions
- Amazon RDS Proxy will simplify relational database usage with Lambda
- Concise logic
- Push orchestration up to AWS Step Functions
- Lambda destinations will simplify and improve asynchronous workflows

## Execution Environment Recap:

- More memory == More CPU and I/O (proportionally)
  - Can also be lower cost
- Use AWS X-Ray to profile your workload
- >1.8 GB memory gets you 2 cores, but you might not use/need it
- Think deeply about your execution model and invocation source needs
  - Not everything needs to be an API
- Thinking async will get you over some of the biggest scaling challenges
- Understand the various aspects to queues, topics, streams and event buses when using them
- Minimize the scope of IAM permissions
  - Leverage tooling like AWS SAM

## Lambda Service Recap:

- Provisioned Concurrency will improve consistency and overall latency of function execution


# Learn Serverless with AWS Training and Certification

Resources created by the experts at AWS to help you learn modern application development



## Free, on-demand courses on serverless, including:

- Introduction to Serverless Development
- Getting in the Serverless Mindset
- AWS Lambda Foundations
- Amazon API Gateway for Serverless Applications
- Amazon DynamoDB for Serverless Architectures



Additional digital and classroom trainings cover modern application development and computing

Visit the Learning Library at <https://aws.training>

# Thank you!

**Chris Munns**

munns@amazon.com

@chrismunns



Please complete the session  
survey in the mobile app.